

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/305748557>

Probabilistic data structures. Part 1. Membership


Presentation · July 2016

DOI: 10.13140/RG.2.1.4836.1203

CITATIONS
0

READS
359

1 author:




Andrii Gakhov


V. N. Karazin Kharkiv National University

17 PUBLICATIONS 4 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:

- 

Time Series Forecasting with Python / MOOC at Udemy [View project](#)
- 

Energy Analysis for software systems [View project](#)



PROBABILISTIC DATA STRUCTURES

ALL YOU WANTED TO KNOW BUT WERE AFRAID TO ASK

PART 1: MEMBERSHIP

Andrii Gakhov
tech talk @ ferret

MEMBERSHIP

Agenda:

- ▶ Bloom Filter
- ▶ Quotient filter

THE PROBLEM

- To determine membership of the element in a large set of elements

BLOOM FILTER

BLOOM FILTER

- proposed by Burton Howard Bloom, 1970
- Bloom filter is a realisation of *probabilistic set* with 3 operations:
 - **add** element into the set
 - test whether an element **is a member** of the set
 - test whether an element **is not a member** of the set
- Bloom filter is described by 2 parameters:
 - **m** - length of the filter
(proportional to expected number of elements **n**)
 - **k** - number of different hash functions
(usually, **k** is much smaller than **m**)
- It doesn't store elements and require about 1 byte per stored data

BLOOM FILTER: ALGORITHM

- Bloom filter is a *bit array* of **m** bits, all set to 0 at the beginning
- **To insert** element into the filter - calculate values of all **k** hash functions for the element and set bit with the corresponding indices
- **To test** if element is in the filter - calculate all **k** hash functions for the element and check bits in all corresponding indices:
 - if all bits are set, then answer is "**maybe**"
 - if at least 1 bit isn't set, then answer is "**definitely not**"
- **Time** needed to insert or test elements is a fixed constant **$O(k)$** , independent from the number of items already in the filter

BLOOM FILTER: EXAMPLE

- Consider Bloom filter of 16 bits ($m=16$)

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

- Consider 2 hash functions ($k=2$):

MurmurHash3 and **Fowler-Noll-Vo**

(to calculate the appropriate index, we divide result by mod 16)

- Add element to the filter: "ferret":

$\text{MurmurHash3}(\text{"ferret"}) = 1, \text{FNV}(\text{"ferret"}) = 11$

0	1	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

BLOOM FILTER: EXAMPLE

- Add element to the filter: "bernau":

$\text{MurmurHash3}(\text{"bernau"}) = 4, \text{FNV}(\text{"bernau"}) = 4$



- Test element: "berlin":

$\text{MurmurHash3}(\text{"berlin"}) = 4, \text{FNV}(\text{"berlin"}) = 12$

Bit **12** is not set, so "berlin" **definitely not in the set**

- Test element: "paris":

$\text{MurmurHash3}(\text{"paris"}) = 11, \text{FNV}(\text{"paris"}) = 4$

Bits **4** and **11** are set, so the element **maybe in the set (false positive)**

BLOOM FILTER: PROPERTIES

- **False positives are possible.** $P(e \in \mathcal{S} | e \notin \mathcal{S}) \approx \left(1 - e^{-\frac{kn}{m}}\right)^k$
(element is not a member, but filter returns like it is a member)
- **False negatives are not possible.** $P(e \notin \mathcal{S} | e \in \mathcal{S}) = 0$
(filter returns that elements isn't a member only if it's not a member)
- Hash functions should be **independent** and **uniformly distributed**. They should also be as **fast as possible**.
(don't use cryptographic, like sha1)
- By choice of **k** and **m** it is possible to decrease false positives probability.
$$k^* = \frac{m}{n} \ln 2$$

BLOOM FILTER: APPLICATIONS

- **Google BigTable, Apache HBase and Apache Cassandra** use Bloom filters to reduce the disk lookups for non-existent rows or columns
- **Medium** uses Bloom filters to avoid recommending articles a user has previously read
- **Google Chrome** web browser used to use a Bloom filter to identify malicious URLs (moved to PrefixSet, Issue 71832)
- The **Squid Web Proxy** Cache uses Bloom filters for cache digests

BLOOM FILTER: PROBLEMS

- The basic Bloom filter **doesn't support deletion**.
- Bloom filters work well when they **fit in main memory**
 - ~1 byte per element (3x-4x times more for Counting Bloom filters, that support deletion)
- What goes wrong when Bloom filters **grow too big to fit in main memory?**
 - On disks with rotating platters and moving heads, Bloom filters choke. A rotational disk performs only 100–200 (random) I/Os per second, and each Bloom filter operation requires multiple I/Os.
 - On flash-based solid-state drives, Bloom filters achieve only hundreds of operations per second in contrast to the order of a million per second in main memory.
- Buffering can help. However, buffering **scales poorly** as the Bloom-filter size increases compared to the in-memory buffer size, resulting in only a few buffered updates per flash page on average.

BLOOM FILTER: VARIANTS

- **Attenuated Bloom filters** use arrays of Bloom filters to store shortest path distance information
- **Spectral Bloom filters** extend the data structure to support estimates of frequencies.
- **Counting Bloom Filters** each entry in the filter instead of a single bit is rather a small counter. Insertions and deletions to the filter increment or decrement the counters respectively.
- **Compressed Bloom filters** can be easily adjusted to the desired tradeoff between size and false-positive rate
- **Bloom Filter Cascade** implements filtering by cascading pipeline of Bloom filters
- **Scalable Bloom Filters** can adapt dynamically to the number of elements stored, while assuring a maximum false positive probability

BLOOM FILTER: PYTHON

- <https://github.com/jaybaird/python-bloomfilter>
pybloom is a module that includes a Bloom Filter data structure along with an implementation of Scalable Bloom Filters
- <https://github.com/seomoz/pyreBloom>
pyreBloom provides Redis backed Bloom Filter using GETBIT and SETBIT

BLOOM FILTER: READ MORE

- http://dmod.eu/deca/ft_gateway.cfm.pdf
- https://en.wikipedia.org/wiki/Bloom_filter
- <https://www.cs.uchicago.edu/~matei/PAPERS/bf.doc>
- <http://gsd.di.uminho.pt/members/cbm/ps/dbloom.pdf>
- <https://www.eecs.harvard.edu/~michaelm/postscripts/im2005b.pdf>

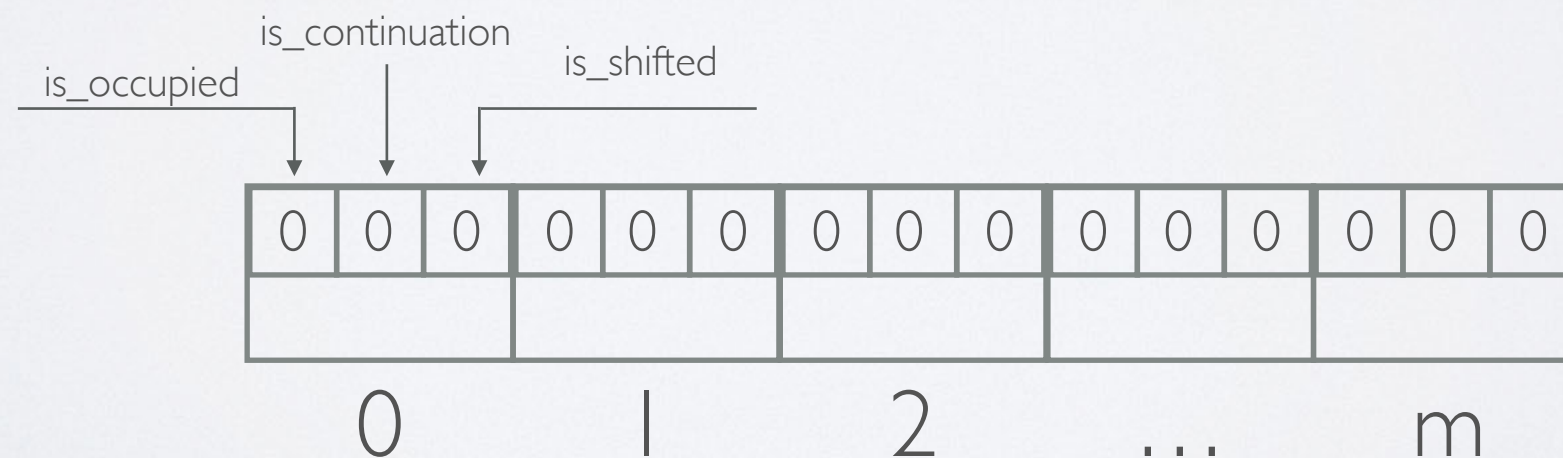
QUOTIENT FILTER

QUOTIENT FILTER

- introduced by Michael Bender et al., 2011
- Quotient filter is a realisation of *probabilistic set* with 4 operations:
 - **add** element into the set
 - **delete** element from the set
 - test whether an element **is a member** of a set
 - test whether an element **is not a member** of a set
- Quotient filter is described by:
 - **p** - size (in bits) for fingerprints
 - **single** hash function that generates fingerprints
- Quotient filter stores **p-bit fingerprints** of elements


QUOTIENT FILTER: ALGORITHM

- Quotient filter is a compact **open hash table** with $m = 2^q$ buckets. The hash table employs *quotienting*, a technique suggested by D. Knuth:
 - the fingerprint f is partitioned into:
 - the r **least significant bits** ($f_r = f \bmod 2^r$, the remainder)
 - the $q = p - r$ **most significant bits** ($f_q = \lfloor f / 2^r \rfloor$, the quotient)
- The remainder is stored in the bucket indexed by the quotient
- Each bucket contains 3 bits, all 0 at the beginning: **is_occupied**, **is_continuation**, **is_shifted**



QUOTIENT FILTER: ALGORITHM

- If two fingerprints \mathbf{f} and \mathbf{f}' have the same quotient ($\mathbf{f}_q = \mathbf{f}'_q$) - it is a **soft collision**. All remainders of fingerprints with the same quotient are stored contiguously in a **run**.
- If necessary, a remainder is **shifted** forward from its original location and stored in a subsequent bucket, wrapping around at the end of the array.
 - **is_occupied** is set when a bucket \mathbf{j} is the canonical bucket ($\mathbf{f}_q = \mathbf{j}$) for some fingerprint \mathbf{f} , stored (somewhere) in the filter
 - **is_continuation** is set when a bucket is occupied but not by the first remainder in a *run*
 - **is_shifted** is set when the remainder in a bucket is not in its canonical bucket




QUOTIENT FILTER: ALGORITHM

- To **test** a fingerprint f :
 - if bucket f_q is not occupied, then the element with fingerprint f definitely not in the filter
 - if bucket f_q is occupied:
 - starting with bucket f_q , scan **left** to locate bucket without set **is_shifted** bit
 - scan **right** with running count (is_occupied: +1, is_continuation: -1) until the running count reaches 0 - when it's the quotient's *run*.
 - compare the remainder in each bucket in the quotient's run with f_r
 - if found, then element *is (probably)* in the filter, else - it is **definitely not** in the filter.
- To **add** a fingerprint f :
 - follow a path similar to test until certain that the fingerprint is definitely not in the filter
 - choose bucket in the current *run* by keeping the sorted order and insert reminder f_r (set is_occupied bit)
 - shift forward all reminders at or after the chosen bucket and update the buckets' bits.

QUOTIENT FILTER: EXAMPLE

- Consider Quotient filter with size $p=8$ and 32-bit signed **MurmurHash3** as h




- Add elements: "amsterdam", "berlin", "london"

$$f_q(\text{"amsterdam"}) = 1, f_q(\text{"berlin"}) = 4, f_q(\text{"london"}) = 7$$

$$f_r(\text{"amsterdam"}) = 164894540, f_r(\text{"berlin"}) = -89622902, f_r(\text{"london"}) = 232552816$$

Insertion at this stage is easy since all canonical slots are not occupied




QUOTIENT FILTER: EXAMPLE

- Add element: "madrid"

$$f_q(\text{"madrid"}) = 1, f_r(\text{"madrid"}) = 249059682$$

The canonical slot 1 is already occupied. The *shifted* and *continuation bits* are not set, so we are at the beginning of the cluster which is also the run's start.

The reminder $f_r(\text{"madrid"})$ is *strongly bigger* than the existing reminder, so it should be shifted right into the next available slot 2 and *shifted/continuation bits* should be set (but not the *occupied bit*, because it pertains to the slot, not the contained reminder).

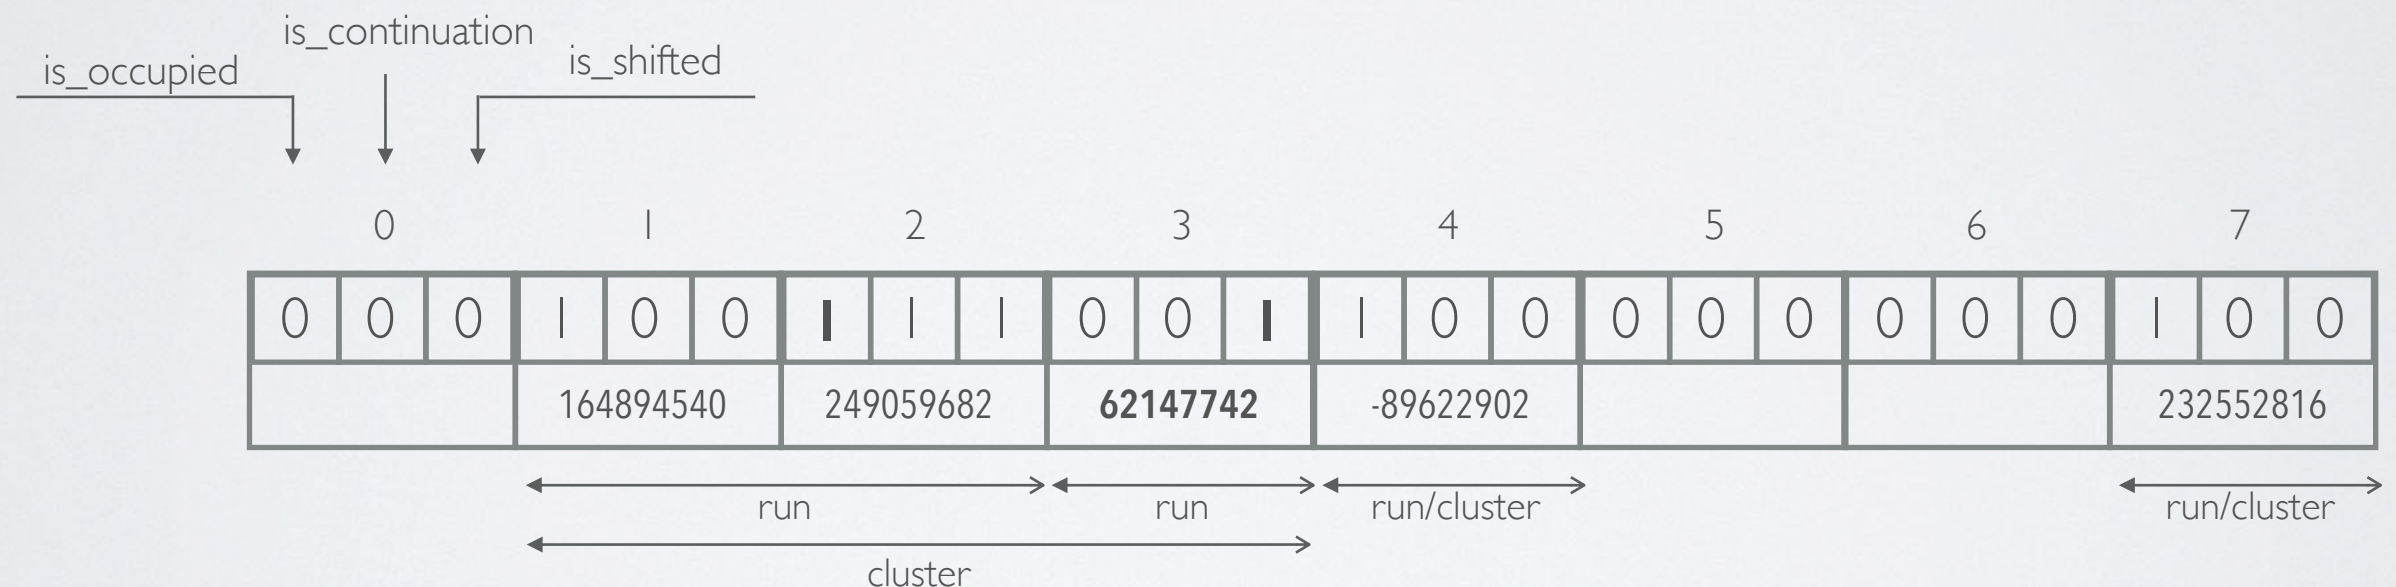


QUOTIENT FILTER: EXAMPLE

- Add element: "ankara"

$$f_q(\text{"ankara"}) = 2, f_r(\text{"ankara"}) = 62147742$$

The canonical slot 2 is not occupied, but already in use. So, the $f_r(\text{"ankara"})$ should be shifted right into the nearest available slot 3 and its *shifted bit* should be set. In addition, we need to flag the canonical slot 2 as occupied by setting the *occupied bit*.



QUOTIENT FILTER: EXAMPLE

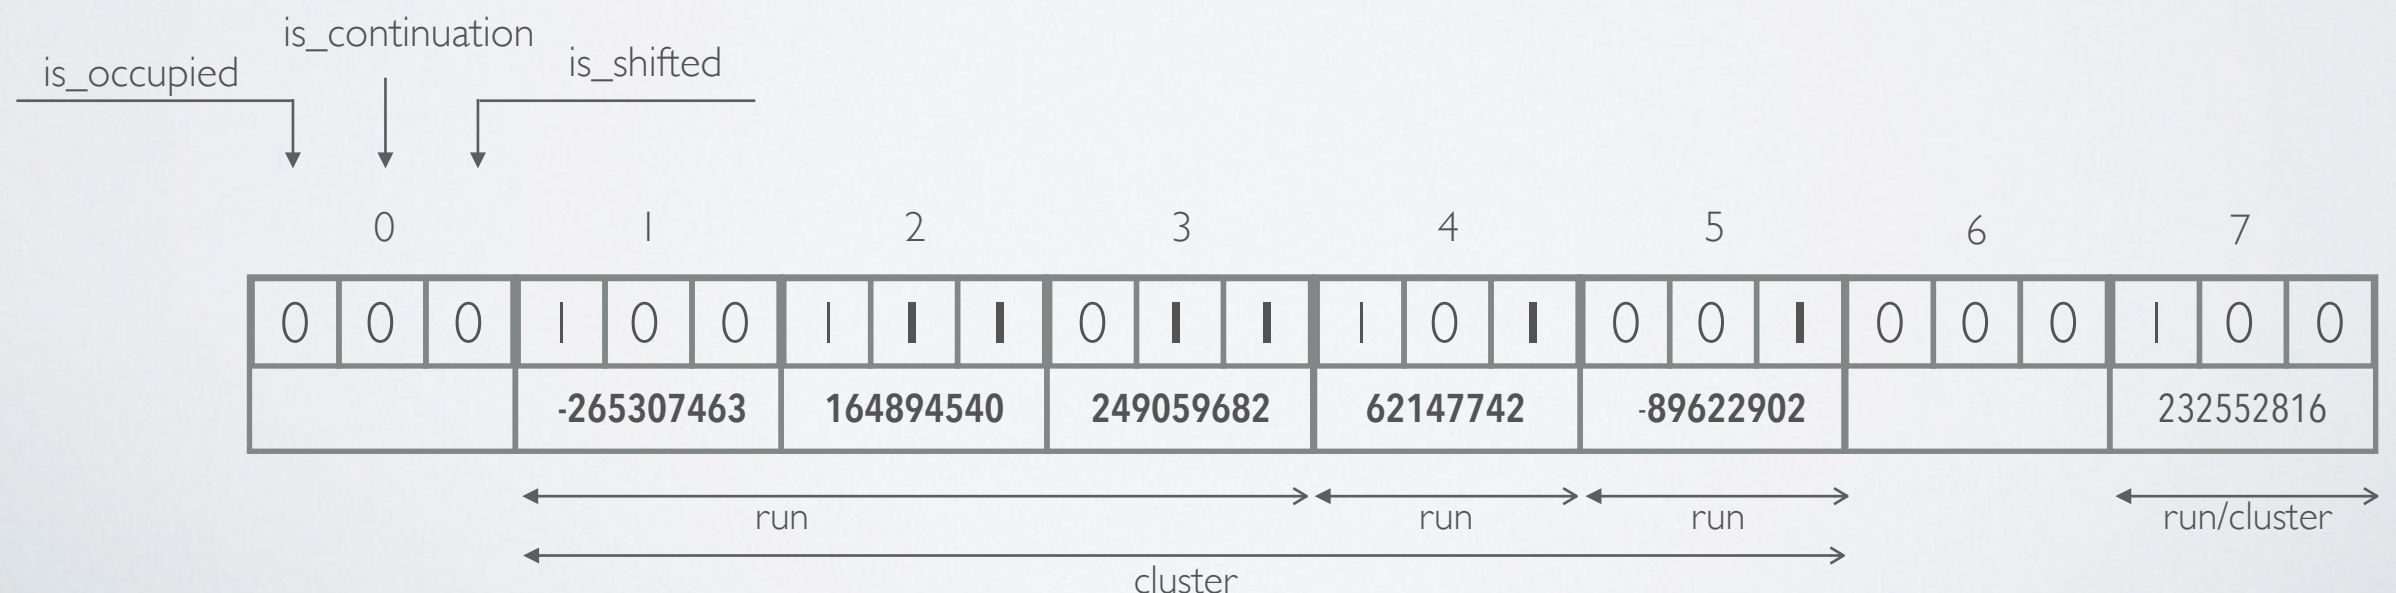
- Add element: "abu dhabi"

$$f_q(\text{"abu dhabi"}) = 1, f_r(\text{"abu dhabi"}) = -265307463$$

The canonical slot 1 is already occupied. The *shifted* and *continuation bits* are not set, so we are at the beginning of the cluster which is also the run's start.

The reminder $f_r(\text{"abu dhabi"})$ is *strongly smaller* than the existing reminder, so all reminders in slot 1 should be shifted right and flagged as continuation and shifted.

If shifting affects reminders from other runs/clusters, we also shift them right and set *shifted bits* (and mirror the continuation bits if they are set there).



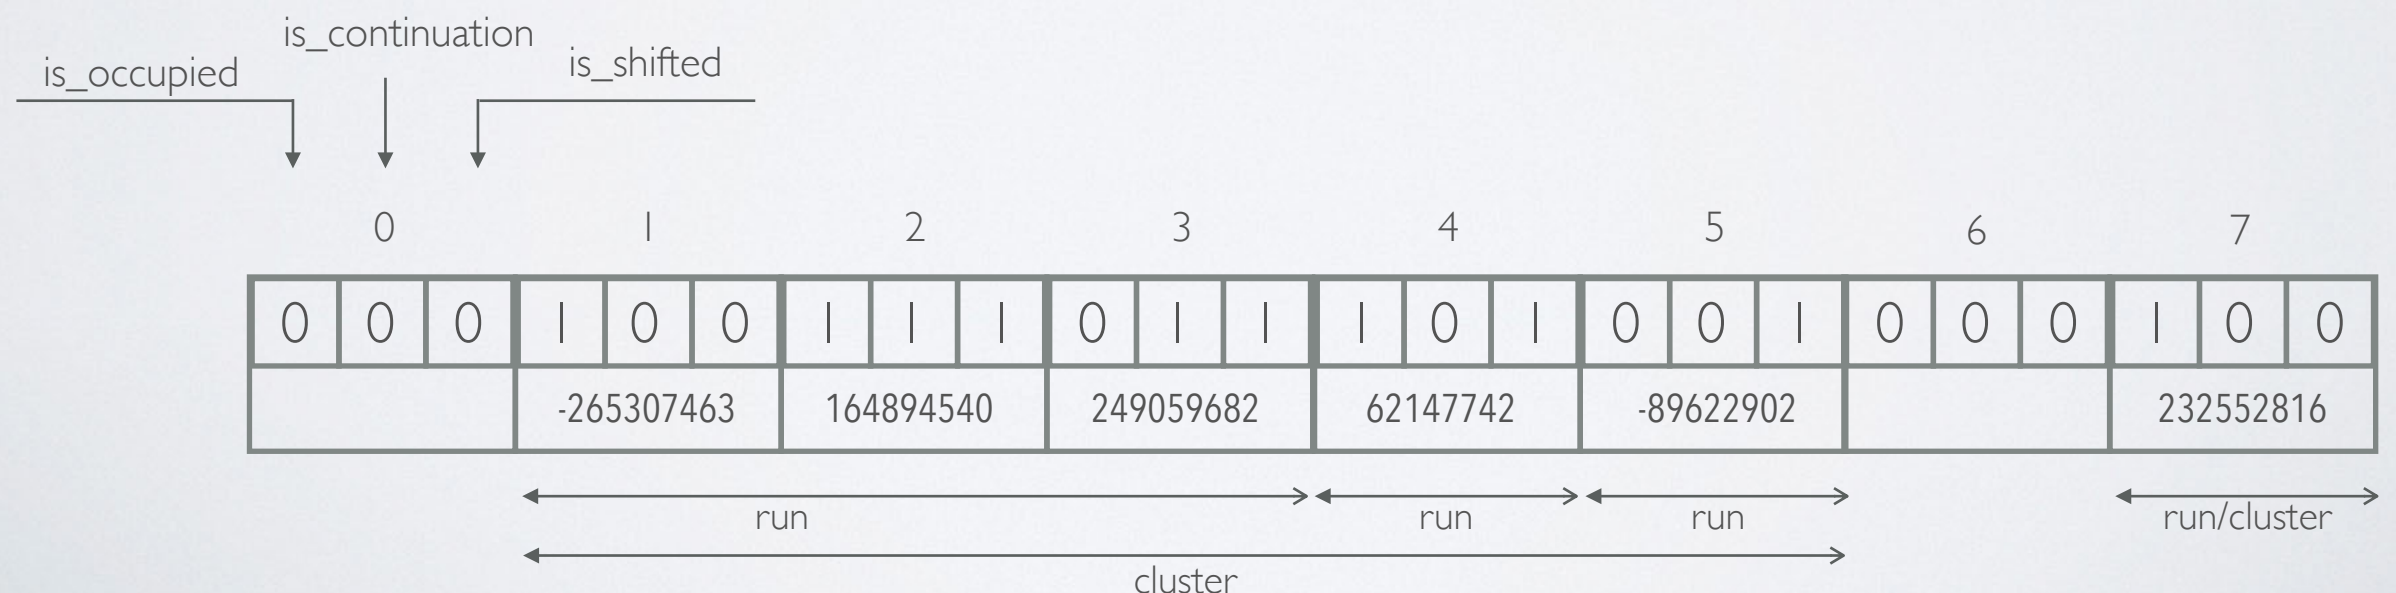
QUOTIENT FILTER: EXAMPLE

- Test element: "ferret"

$$f_q(\text{"ferret"}) = 1, f_r(\text{"ferret"}) = 122150710$$

The canonical slot 1 is already occupied and it's the start of the run. Iterate through the run and compare $f_r(\text{"ferret"})$ with existing reminders until we found the match, found reminder that is strongly bigger, or hit the run's end.

We start from slot 1 which reminder is smaller, so we continue to slot 2. Reminder in the slot 2 is already bigger than $f_r(\text{"ferret"})$, so we conclude that "ferret" **definitely not in the filter**.



QUOTIENT FILTER: EXAMPLE

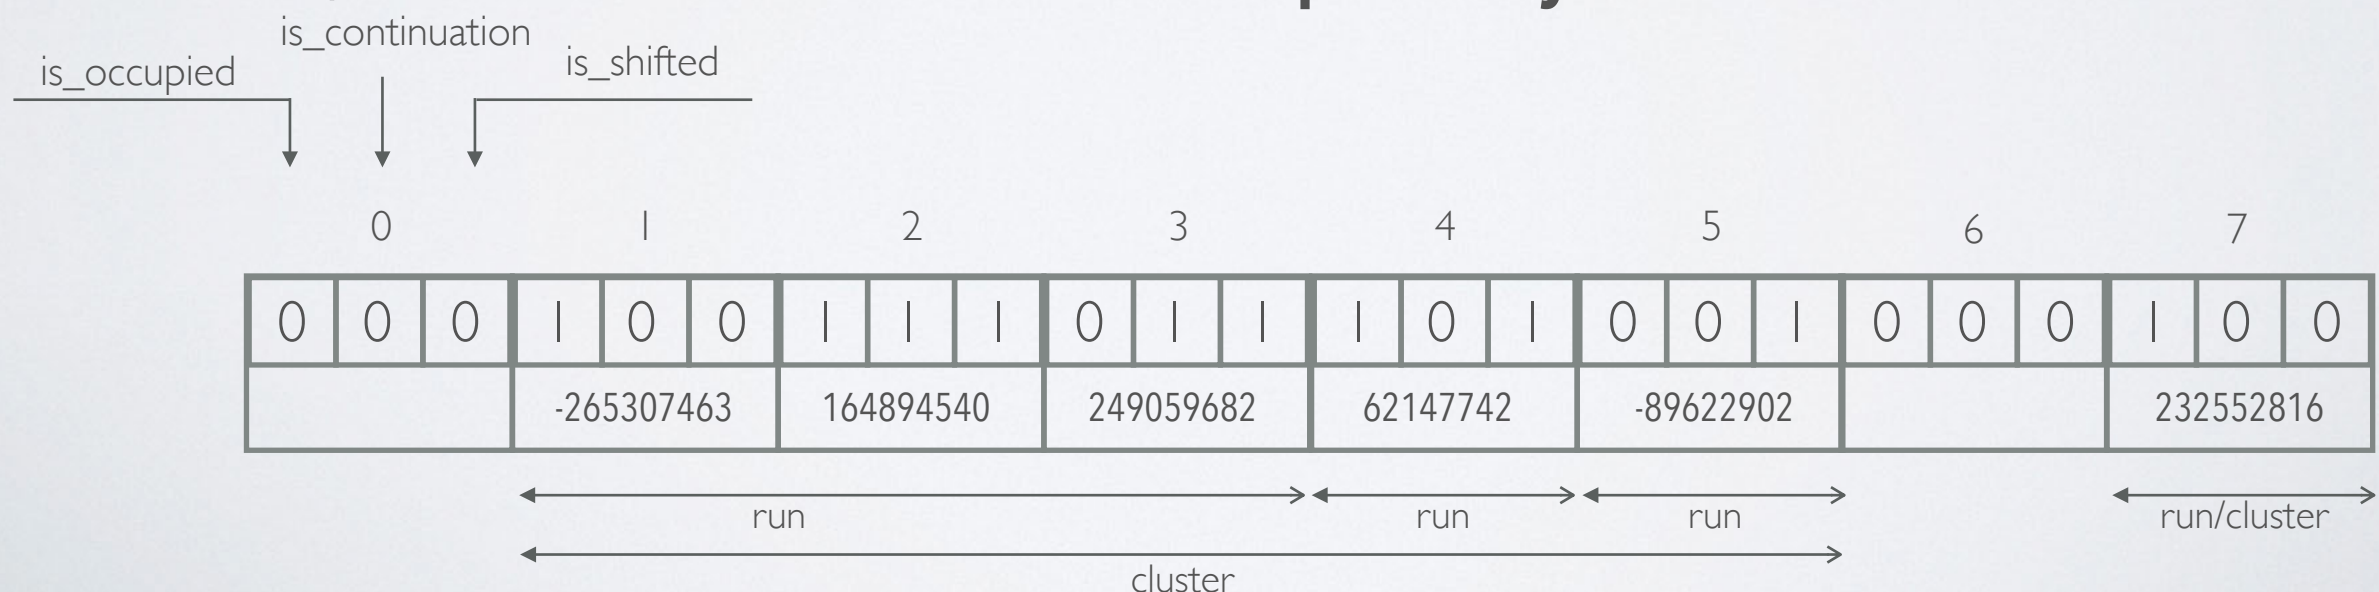
- Test element: "berlin"

$$f_q(\text{"berlin"}) = 4, f_r(\text{"berlin"}) = -89622902$$

The canonical slot 4 is already occupied, but shifted bit is set, so the run for which it is canonical slot exists, but is shifted right.

First, we need to find a run corresponding to the canonical slot 4 in the current cluster., so we scan left and count occupied slots. There are 2 occupied slots found (indices 1 and 2), therefore our run is the 3rd in the cluster and we can scan right until we found it (count slots with not set *continuation bit*).

Our run starts in the slot 5 and we start comparing $f_r(\text{"berlin"})$ with existing values and found match, so we can conclude that "berlin" **probably in the filter**.



QUOTIENT FILTER: PROPERTIES

- **False positives** are **possible**.

$$P(e \in \mathcal{S} \mid e \notin \mathcal{S}) \leq \frac{1}{2^r}$$

(element is not a member, but filter returns like it is a member)

- **False negatives** are **not possible**.

$$P(e \notin \mathcal{S} \mid e \in \mathcal{S}) = 0$$

(filter returns that elements isn't a member only if it's not a member)

- Hash function should generate **uniformly distributed** fingerprints.
- The length of most runs is $O(1)$ and it is highly likely that all runs have length $O(\log m)$
- Quotient filter efficient for large number of elements (~ 1 B for 64-bit hash function)

QUOTIENT FILTER VS BLOOM FILTER

- Quotient filters are about 20% **bigger** than Bloom filters, but **faster** because each access requires evaluating only a **single hash function**
- Results of comparison of in-RAM performance (M. Bender et al.):
 - **inserts**: BF: 690 000 inserts per second, QF: 2 400 000 insert per second
 - **lookups**: BF: 1 900 000 lookups per second, QF: 2 000 000 lookups per second
- Lookups in Quotient filters incur a single cache miss, as opposed to at least two in expectation for a Bloom filter.
- Two Quotient Filters can be efficiently **merged without affecting their false positive** rates. This is not possible with Bloom filters.
- Quotient filters support deletion

QUOTIENT FILTER: IMPLEMENTATIONS

- <https://github.com/vedantk/quotient-filter>
Quotient filter in-memory implementation written in C
- <https://github.com/dsx724/php-quotient-filter>
Quotient filter implementation in pure PHP
- <https://github.com/bucaojit/QuotientFilter>
Quotient filter implementation in Java

QUOTIENT FILTER: READ MORE

- http://static.usenix.org/events/hotstorage11/tech/final_files/Bender.pdf
- <http://arxiv.org/pdf/1208.0290.pdf>
- https://en.wikipedia.org/wiki/Quotient_filter
- <http://www.vldb.org/pvldb/vol6/p589-dutta.pdf>

THANK YOU

- ▶ @gakhov
- ▶ [linkedin.com/in/gakhov](https://www.linkedin.com/in/gakhov)
- ▶ www.datacrucis.com