

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/305748721>

# Probabilistic data structures. Part 2. Cardinality

Presentation · July 2016

DOI: 10.13140/RG.2.1.2476.8244

CITATIONS  
0

READS  
204

1 author:





Andrii Gakhov  
V. N. Karazin Kharkiv National University

17 PUBLICATIONS 4 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:

 Ph.D. in Mathematical Modelling and Numerical Methods [View project](#)

 Energy Analysis for software systems [View project](#)



# PROBABILISTIC DATA STRUCTURES

ALL YOU WANTED TO KNOW BUT WERE AFRAID TO ASK

## PART 2: CARDINALITY

Andrii Gakhov  
tech talk @ ferret

# CARDINALITY

## Agenda:

- ▶ Linear Counting
- ▶ LogLog, SuperLogLog, HyperLogLog, HyperLogLog++

# THE PROBLEM

- To determine the *number of distinct elements*, also called the **cardinality**, of a large set of elements where duplicates are present

Calculating the exact cardinality of a multiset requires an amount of memory proportional to the cardinality, which is impractical for very large data sets.

# LINEAR COUNTING

# LINEAR COUNTING: ALGORITHM

- Linear counter is a **bit map** (hash table) of size **m** (all elements set to 0 at the beginning).
- Algorithm consists of a few steps:
  - for every element calculate hash function and set the appropriate bit to 1
  - calculate the fraction **V** of empty bits in the structure (divide the number of empty bits by the bit map size **m** )
  - estimate cardinality as  **$n \approx -m \ln V$**



# LINEAR COUNTING: EXAMPLE

- Consider linear counter with 16 bits ( $m=16$ )

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

- Consider **MurmurHash3** as the hash function **h**  
(to calculate the appropriate index, we divide result by mod 16)
- Set of 10 elements: "*bernau*", "*bernau*", "*bernau*", "*berlin*", "*kiev*", "*kiev*", "*new york*", "*germany*", "*ukraine*", "*europa*" (NOTE: the real cardinality  $n = 7$ )

$h(\text{"bernau"}) = 4$ ,  $h(\text{"berlin"}) = 4$ ,  $h(\text{"kiev"}) = 6$ ,  $h(\text{"new york"}) = 6$ ,  
 $h(\text{"germany"}) = 14$ ,  $h(\text{"ukraine"}) = 7$ ,  $h(\text{"europa"}) = 9$

0	0	0	0	1	0	1	1	0	1	0	0	0	0	1	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

# LINEAR COUNTING: EXAMPLE

0	0	0	0	1	0	1	1	0	1	0	0	0	0	1	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

number of empty bits: 11

$$m = 16$$

$$V = 11 / 16 = 0.6875$$

- Cardinality estimation is

$$n \approx -16 * \ln(0.6875) = 5.995$$



# LINEAR COUNTING: READ MORE

- [http://dblab.kaist.ac.kr/Prof/pdf/Whang1990\(linear\).pdf](http://dblab.kaist.ac.kr/Prof/pdf/Whang1990(linear).pdf)
- [http://www.codeproject.com/Articles/569718/  
CardinalityplusEstimationplusinplusLinearplusTimep](http://www.codeproject.com/Articles/569718/CardinalityplusEstimationplusinplusLinearplusTimeplusSpaceplusComplexityplusAnalysisplusImplementationplus)

# **HYPERLOGLOG**

# HYPERLOGLOG: INTUITION

- The cardinality of a multiset of *uniformly distributed numbers* can be estimated by the maximum number of leading zeros in the binary representation of each number. If such value is **k**, then the number of distinct elements in the set is  $2^k$

$f =$

0	1	2	3												
0	0	1	0	0	0	1	0	0	0	1	0	0	0	0	0

←→  
leading zeros

**rank** = number of leading zeros + 1, e.g.  $\text{rank}(f) = 3$

$P(\text{rank}=1) = 1/2$  - probability to find a binary representation, that starts with 1

$P(\text{rank} = 2) = 1/2^2$  - probability to find a binary representation, that start with 01

...

$P(\text{rank}=k) = 1/2^k$

- Therefore, for  $2^k$  binary representations we shell find at least one representation with **rank = k**
- If we remember the maximal rank we've seen and it's equal to **k**, then we can use  $2^k$  as the approximation of the number of elements

# HYPERLOGLOG

- proposed by Flajolet et. al., 2007
- an extension of the Flajolet–Martin algorithm (1985)
- HyperLogLog is described by 2 parameters:
  - **p** – number of bits that determine a bucket to use averaging ( $m = 2^p$  is the number of buckets/substreams)
  - **h** - hash function, that produces *uniform* hash values
- The HyperLogLog algorithm is able to estimate cardinalities of  $> 10^9$  with a typical error rate of 2%, using 1.5kB of memory (Flajolet, P. et al., 2007).



# HYPERLOGLOG: ALGORITHM

- HyperLogLog uses **randomization** to approximate the cardinality of a multiset. This randomization is achieved by using hash function **h**
- Observe the maximum **number of leading zeros** that for all hash values:
  - If the bit pattern  **$0^{L-1} 1$**  is observed at the beginning of a hash value (so,  $\text{rank} = L$ ), then a *good estimation of the size* of the multiset is  **$2^L$** .



# HYPERLOGLOG: ALGORITHM

- **Stochastic averaging** is used to reduce the large variability:
  - The input stream of data elements  $S$  is divided into  $m$  substreams  $S_i$  using the first  $p$  bits of the hash values ( $m = 2^p$ ).
  - In each substream, the rank (after the initial  $p$  bits that are used for substreaming) is measured independently.
  - These numbers are kept in an array of registers  $M$ , where  $M[i]$  stores the maximum rank it seen for the substream with index  $i$ .
- The **cardinality estimation** is calculated computes as the normalized bias corrected harmonic mean of the estimations on the substreams

$$DV_{HLL} = \text{const}(m) \cdot m^2 \cdot \left( \sum_{j=1}^m 2^{-M_j} \right)^{-1}$$

# HYPERLOGLOG: EXAMPLE

- Let's use **p=3** bits to define a bucket (then  $m=2^3=8$  buckets).

	0	1	2	3	4	5	6	7
<b>M</b>	0	0	0	0	0	0	0	0

- Consider **L=8** bits hash function **h**
- Index elements "berlin" and "ferret":

$h(\text{"berlin"}) = \mathbf{011}0111$

$h(\text{"ferret"}) = \mathbf{110}0011$

- Define buckets and calculate values to store:

(use first  $p=3$  bits for buckets and least  $L-p=5$  bits for ranks)

- bucket("berlin")** = 011 = 3    **value("berlin")** = rank(0111) = 2
- bucket("ferret")** = 110 = 6    **value("ferret")** = rank(0011) = 3

	0	1	2	3	4	5	6	7
<b>M</b>	0	0	0	<b>2</b>	0	0	<b>3</b>	0

# HYPERLOGLOG: EXAMPLE

- Index element "kharkov":
  - $h(\text{"kharkov"}) = 1100001$
  - $\text{bucket}(\text{"kharkov"}) = 110 = 6$     $\text{value}(\text{"kharkov"}) = \text{rank}(0001) = 4$
  - $M[6] = \max(M[6], 4) = \max(3, 4) = 4$

	0	1	2	3	4	5	6	7
<b>M</b>	0	0	0	<b>2</b>	0	0	<b>4</b>	0

- Estimate the cardinality by the HLL formula ( $C \approx 0.66$ ):
$$DV_{HLL} \approx 0.66 * 8^2 / (2^{-2} + 2^{-4}) = 0.66 * 204.8 \approx 135 \neq 3$$

**NOTE: For small cardinalities HLL has a strong bias!!!**

# HYPERLOGLOG: PROPERTIES

- **Memory requirement doesn't grow linearly with  $L$**  (unlike MinCount or Linear Counting) - for hash function of  $L$  bits and precision  $p$ , required memory:

$$\lceil \log_2 (L + 1 - p) \rceil \cdot 2^p \text{ bits}$$

- original HyperLogLog uses 32 bit hash codes, which requires  **$5 \cdot 2^p$  bits**
- **It's not necessary to calculate the full hash code** for the element
  - first  $p$  bits and number of leading zeros of the remaining bits are enough
- There are no evidence that some of popular hash functions (MD5, Sha1, Sha256, Murmur3) performs significantly better than others.



# HYPERLOGLOG: PROPERTIES

- Algorithm has large error for small cardinalities.
  - For instance, for  $n = 0$  the algorithm always returns roughly **0.7m**
  - To achieve better estimates for small cardinalities, use *LinearCounting* below a threshold of **5m/2**
- The **standard error** can be estimated as:

$$\sigma = \frac{1.04}{\sqrt{2^p}}$$

so, if we use 16 bits ( $p=16$ ) for bucket indices, we receive the standard error in **0.40625%**



# HYPERLOGLOG: APPLICATIONS

- *PFCOUNT* in **Redis** returns the approximated cardinality computed by the HyperLogLog data structure (<http://antirez.com/news/75>)
- Redis implementation uses 12Kb per key to count with a standard error of 0.81%, and there is no limit to the number of items you can count, unless you approach  $2^{64}$  items

# HYPERLOGLOG: READ MORE

- <http://algo.inria.fr/flajolet/Publications/DuFl03-LNCS.pdf>
- <http://algo.inria.fr/flajolet/Publications/FlFuGaMe07.pdf>
- <https://stefanheule.com/papers/edbt13-hyperloglog.pdf>
- <https://highlyscalable.wordpress.com/2012/05/01/probabilistic-structures-web-analytics-data-mining/>
- [https://hal.archives-ouvertes.fr/file/index/docid/465313/filename/sliding\\_HyperLogLog.pdf](https://hal.archives-ouvertes.fr/file/index/docid/465313/filename/sliding_HyperLogLog.pdf)
- <http://stackoverflow.com/questions/12327004/how-does-the-hyperloglog-algorithm-work>

**HYPERLOGLOG++**

# HYPERLOGLOG++

- proposed by Stefan Heule et. al., 2013 for Google PowerDrill
- an improved version of HyperLogLog (Flajolet et. al., 2007)
- HyperLogLog++ is described by 2 parameters:
  - **p** – number of bits that determine a bucket to use averaging ( $m = 2^p$  is the number of buckets/substreams)
  - **h** - hash function, that produces *uniform* hash values
- The HyperLogLog++ algorithm is able to estimate cardinalities of  $\sim 7.9 \cdot 10^9$  with a typical error rate of 1.625%, using 2.56KB of memory (Micha Gorelick and Ian Ozsvald, High Performance Python, 2014).



# HYPERLOGLOG++: IMPROVEMENTS

- **use 64-bit hash function**
  - algorithm that only uses the hash code of the input values is limited by the number of bits of the hash codes when it comes to accurately estimating large cardinalities
  - In particular, a hash function of  $L$  bits can at most distinguish  $2^L$  different values, and as the cardinality  $n$  approaches  $2^L$  hash collisions become more and more likely and accurate estimation gets impossible
    - if the cardinality approaches  $2^{64} \approx 1.8 \cdot 10^{19}$ , hash collisions become a problem
- **bias correction**
  - original algorithm overestimates the real cardinality for small sets, but most of the error is due to bias.
- **storage efficiency**
  - uses different encoding strategies for hash values, variable length encoding for integers, difference encoding



# HYPERLOGLOG++ VS HYPERLOGLOG

- **accuracy is significantly better** for large range of cardinalities and equally good on the rest
- sparse representation allows for a **more adaptive use of memory**
- if the cardinality **n** is much smaller than **m**, then HyperLogLog++ **requires significantly less memory**
- For cardinalities between 12000 and 61000, the bias correction allows for a **lower error** and avoids a spike in the error when switching between sub-algorithms.
- 64 bit hash codes allow the algorithm to estimate cardinalities well beyond **1 billion**

# HYPERLOGLOG++: APPLICATIONS

- *cardinality* metric in **Elasticsearch** is based on the HyperLogLog++ algorithm for big cardinalities (adaptive counting)
- **Apache DataFu**, collection of libraries for working with large-scale data in Hadoop, has an implementation of HyperLogLog++ algorithm

# HYPERLOGLOG++: READ MORE

- <http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/40671.pdf>
- <https://research.neustar.biz/2013/01/24/hyperloglog-googles-take-on-engineering-hll/>

# THANK YOU

- ▶ @gakhov
- ▶ [linkedin.com/in/gakhov](https://www.linkedin.com/in/gakhov)
- ▶ [www.datacrucis.com](http://www.datacrucis.com)